

Contents

1	Desiderata and Examples	2
1.1	Some examples	3
1.2	Mixed strategies	8
2	DS and IDDS	10
2.1	Basic Definitions	11
2.2	DS	12
2.3	IDDS	13
3	Nash equilibrium: Pure Strategies	16
3.1	Equal payoff property	16
3.2	Definition and existence	16
4	Nash Equilibrium: Mixed	19
4.1	Payoffs for mixed strategies.	19
4.2	Best-response diagram	23
4.3	Mixed NE as maximum expected payoff	25
5	Refinements, QRE and CE	27
5.1	Payoff dominance vs risk dominance.	27
5.2	Correlated equilibrium	30
6	Behavioral Considerations.	31

Chapter 4: Solving Normal Form Games

There are several different solution concepts for NFGs. Before presenting them, we first list nice properties that we might hope for in a solution, and run through several examples to build intuition. Next we present definitions on which the concepts are built, and define the first two solution concepts, DS and IDDS. The foremost solution concept, NE is presented in two parts, pure and mixed. Then we briefly mention other concepts, such as CE and QRE, that are sometimes useful.

[note for next edition: don't forget to use commented out HD material]

1 Desiderata and Examples

What should it mean to solve a game? What might we expect a solution to do?

Here is a list of some desirable properties.

A solution of a game should:

1. Be consistent with rationality (SEU, EUH).
2. Always exist.
3. Be unique.
4. Be efficient: either (a) maximize average or total payoff (“transferrable” utility), or perhaps just (b) be Pareto optimal.
5. Be “reasonable”: sensible players could be persuaded to stick with it.
6. Have empirical support, and predict accurately what actual human players do in the lab or field.

Unfortunately, we will see that there is no solution concept that can satisfy all these desiderata. The leading candidate, Nash equilibrium, always satisfies the first two, but

fairly simple games can be found where it fails to satisfy each of the others. Other solution concepts do better on some desiderata and worse on others.

1.1 Some examples

[[half class has Row role, other half has Column. Each person chooses what to do, votes are tallied, explanations elicited.]]

Game 1: Pigs

	L	D
l	1,4	-1,6
d	3,2	0,0

Game 2: IDDSb

	t_1	t_2	t_3
s_1	4, 3	2, <u>7</u>	<u>3</u> , 4
s_2	<u>5</u> , <u>5</u>	<u>5</u> , -1	-4, 2

- Convention is that row player is called player 1 and column player is called player 2. Best-responses are underlined.
- Row has 2 possible strategies: s_1 and s_2 .
- Column has 3 possible strategies: t_1 , t_2 , and t_3 .
- One way to think about solving games is to see how each player responds best to each possible situation.
- Iteration (i)
 - Looking at Row player in Game 2, we can state the best-responses as:

if t_1 then s_2 is best-response.

if t_2 then s_2 is best-response.

if t_3 then s_1 is best-response.

– Thus, initially we can not eliminate any rows.

• Iteration (ii)

– Then look at Columns' best-responses:

if s_1 then t_2 is best-response.

if s_2 then t_1 is best-response.

– Thus, t_3 is never a best-response.

* If s_1 , then t_3 is dominated by t_2 .

* If s_2 , then t_3 is dominated by t_1 .

* Since t_3 is not a best-response to either pure strategy, it is possible that t_3 is dominated by some mixture of t_1 and t_2 .

* Consider, for example, playing t_1 and t_2 with probability $p(t_1) = 0.6$ and $p(t_2) = 0.4$:

· against s_1 : $\pi_2(0.6t_1 + 0.4t_2) = \frac{18+28}{10} = 4.6$ better than payoff 4 from t_3 .

· against s_2 : $\pi_2(0.6t_1 + 0.4t_2) = \frac{30+4}{10} = 2.6$ better than payoff 2 from t_3 .

* you can check that any such mixed strategy with $0.5 < p(t_1) < 0.75, p(t_2) = 1 - p(t_1)$ dominates t_3 .

– We can eliminate t_3 and obtain the reduced game:

	t_1	t_2
s_1	4, 3	2, 7
s_2	5, 5	5, -1

- Iteration (iii)

- Then look at row's best-responses:

if t_1 then s_2 is best-response.

if t_2 then s_2 is best-response.

- – Thus, now s_2 dominates s_1 and we can delete s_1 to obtain the reduced game:

	t_1	t_2
s_2	5, 5	5, -1

- Iteration (iv).

- Now, column is facing only s_2 , and clearly t_1 dominates t_2 ,

- Thus we can eliminate t_2 to obtain the solution (s_2, t_1) :

	t_1
s_2	5, 5

- Now follow the same procedure with Game 1.

- The Row player (the “little pig”) has a dominated strategy, l (“push the food lever”).

- Assuming that Row always plays d (“wait by the food dispenser”), the Column player (“big pig”) best responds by playing L.

- The solution (d,L) has payoff (3,2), illustrating that in game theory as in Scripture, things may work out well for the meek!

- We will see later that this solution concept is called iterated deletion of dominated strategies (IDDS) and that it is a special case of Nash equilibrium (NE).
- The IDDS procedure worked for the last 2 games, but there are equally simple games where it doesn't get us very far.

Game 3: Driving

	L [p]	R [$1 - p$]
L	1,1	-10,-10
R	-10,-10	1,1

- No strategy is dominated in Game 3 — you want to make the same choice as the other players as to which side of the road to drive on.
- Population game interpretation: There are a large number of other drivers you may encounter, and fraction p of them choose L.
 - Your expected payoff is $p - 10(1 - p) = 11p - 10$ if you play L, and $-10p + 1(1 - p) = 1 - 11p$ if R.
 - You are better off choosing L if the expected payoff advantage to L ($11p - 10 - (1 - 11p) = 22p - 11$) is positive, i.e., if $p > 0.5$
 - Likewise, you are better off choosing R if more than half of other drivers choose R, i.e., if $1 - p > 0.5$.
 - If $p = 0.5$, as it seems to be in some cities you may have visited, then R and L are equally bad.
 - Indeed, if you could choose to stay home and get payoff 0, you'd prefer that if $p \in [\frac{1}{11}, \frac{10}{11}]$.
 - Thus we have multiple solutions. If everyone else plays R then you'd also prefer R, but if everyone else plays L then you'd prefer that. We will see that these

solutions are examples of pure NE. And the $p = 0.5$ bad solution is an example of a mixed NE.

This sort of situation is called a **coordination game**.

- Basic idea is that there is a payoff advantage to being on the “same page” as the other players.
- There are various sorts of coordination games, as we will see. Consider one famous example, known to economists as “Battle of the Sexes”:
 - Two players, each with two possible actions.
 - traditionally, a heterosexual couple choosing Boxing Match or Opera, but you can make it more interesting if you like. The idea is that they would rather be together, but each has a preferred event.

Game 4: Battle of the Sexes

	t_1	t_2
s_1	0, 0	<u>20</u> , <u>40</u>
s_2	<u>40</u> , <u>20</u>	0, 0

- Looking at the best-responses, starting with Row:
 - if t_1 then s_2 is best-response.
 - if t_2 then s_1 is best-response.
- Thus, row does not have any dominated strategies.
- Column’s best-responses:
 - if s_1 then t_2 is best-response.
 - if s_2 then t_1 is best-response

- Thus, column does not have any dominated strategies either and the IDDS procedure has no traction.
- Another solution idea is that nobody wants to change unilaterally, as in Nash equilibrium (as we will see). We have 2 NE in pure strategies $NE = \{(s_2, t_1), (s_1, t_2)\}$ (and a third lurking in mixed strategies, to be discussed below).
- Note that player 1 prefers (s_2, t_1) and player 2 prefers (s_1, t_2) . We have two predictions, at least. In games in the lab with distinct populations for row players and column players, typically one of these two will emerge eventually as the convention that almost everyone follows.
- There are also other sorts of coordination games we will see later where there is a different tension between predicted NE (risk-dominant vs. payoff dominant).

1.2 Mixed strategies

[[maybe move this to the previous chapter?]]

A **mixed strategy** is a probability distribution over the set of pure strategies.

There are two interpretations of mixed strategies. Both have their uses.

1. each pure strategy is followed by a given fraction of a large population, possibly 0. The fractions are non-negative and sum to 1.0.
2. There is only one player in each role (e.g., Row or Column), and those players make choices as if flipping a coin, spinning a roulette wheel, or rolling dice so as to pick each pure strategy with a given probability.

In some games, there is an advantage to being unpredictable,

- (American) football example: running 30% of the time on first down.
- (soccer) football example: penalty kicks (L,M,R)

Example:

Game 5: Conflict

	t_1	t_2
s_1	$\underline{3}, 1$	$1, \underline{3}$
s_2	$0, \underline{5}$	$\underline{4}, 2$

- Look at the best-responses, first for Row.

s_1 is best-response to t_1 .

s_2 is best-response to t_2 .

- Similarly, looking at columns' best-responses:

t_2 is best-response to s_1 .

t_1 is best-response to s_2 .

- No dominated strategies, no NE in pure strategies. IDDS is no help.

So we need to find a mixed NE.

This seems like a good juncture to discuss mixed strategies in more detail.

- The set of pure strategies is $S = \{s_1, s_2\}$ for the case of two pure strategies.
- $\Delta(S)$ denotes the set of all possible mixed strategies. The set of mixed strategies $\Delta(S)$ is a simplex:

$$\Delta(S) = \{ps_1 + (1-p)s_2\} \text{ where } p \in [0, 1]$$

- If $p = 0$ then pure strategy s_2 is obtained. This means that the set of pure strategies is a proper subset of the set of mixed strategies: $S \subset \Delta(S)$.

- For three pure strategies then $\Delta(S)$ is a three dimensional triangle.

$$\Delta(S) = \{ps_1 + qs_2 + (1-p-q)s_3\} \text{ where } p, q \in [0, 1] \text{ and } p+q \in [0, 1]$$




Figure 1: When a player has three pure strategies, the set of mixed strategies is the equilateral triangle with vertices $(1, 0, 0)$, $(0, 1, 0)$, $(0, 0, 1)$.

- For n possible pure strategies the simplex is a hypertriangle.
- So, a mixed strategy is a probability distribution across the set of pure strategies.
 - Note that the realization of a mixed strategy is one of the pure strategies.
- Picking the mixing probabilities p and q is the trick.
- We will see that even games (as above) with no pure strategy NE still have a NE in mixed strategies.

2 DS and IDDS

With these examples in mind, we can introduce the ideas more formally.

2.1 Basic Definitions

A pure strategy $s_i \in S_i$ is **weakly dominant** if:

$$f_i(s_i, s_{-i}) \geq f_i(s'_i, s_{-i}) \quad \forall \left\{ \begin{array}{l} s'_i \in S_i \\ s_{-i} \in S_{-i} \end{array} \right\} \quad (1)$$

Thus, a weakly dominant strategy is a best-response to *all* possible strategy profiles of the other players.

- Fix what the other players are doing and then choose your best strategy.
- Then vary what everyone else is doing, and if that strategy always is still a best response, then it is weakly dominant.

A pure strategy is said to be **strictly dominant** if the inequality in equation (1) is strict for all other pure strategies $s'_i \neq s_i \in S_i$.

Of course, most games that you will encounter will not have dominant strategies — only in special situations does a player have a single strategy that will always work well. The opposite case is more common, when some players have some strategies that they would never want to use.

A pure strategy $s_i \in S_i$ is **weakly dominated** if:

$$\exists \quad s'_i \in S_i \quad \text{such that} \quad f_i(s_i, s_{-i}) \leq f_i(s'_i, s_{-i}) \quad \forall s_{-i} \in S_{-i} \quad (2)$$

- Thus, a strategy is weakly dominated if there exists a different pure strategy that always yields at least as high a payoff.
- A strategy is **strictly dominated** if the inequality in (2) is strict.
- A pure strategy $s_i \in S_i$ can be dominated by a mixed strategy x , if

$$\exists \quad x \in \Delta(S_i) \quad \text{such that} \quad f_i(s_i, s_{-i}) \leq f_i(x, s_{-i}) \quad \forall s_{-i} \in S_{-i} \quad (3)$$

- Like our Game 2 example: $0.6t_1 + 0.4t_2$ dominates t_3 .

Now we are ready for the key definition.

A pure strategy $s'_i \in S_i$ is a **best-response** to s_{-i} if:

$$f_i(s'_i, s_{-i}) \geq f_i(s_i, s_{-i}) \quad \forall s_i \in S_i \quad (4)$$

If equation (4) holds, we write $s'_i \in B_i(s_{-i})$. So $B_i(s_{-i})$ is the set of all pure strategy best responses that player i has to a particular strategy profile by the other players, s_{-i} . If there is just one such strategy, we sometimes (a bit carelessly) write $s'_i = B_i(s_{-i})$.

2.2 DS

If all players have a strictly dominant strategy, then the game is said to have a **dominant strategy (DS)** solution.

DS definitely satisfy desiderata 1) [consistent with rationality] and 3) [unique]. DS also typically satisfy 5) [reasonable], although this might be arguable in a few cases. Of course, desideratum 2) [existence] fails badly for DS since dominant strategies are a bit special for any single player, much less for all players.

Prisoner's dilemma is the most famous game with a DS solution.

- Original story: s_1 and t_1 are remain silent, s_2 and t_2 are confess. (Other stories concern arms races, cartel behavior, etc. etc.)
- The original payoffs are negative (years in jail) but since even cardinal utility is defined up to a positive affine transformation (adding a constant in this case), we can make the payoffs positive and still represent the same set of preferences. We'll return later to the question of cardinal vs ordinal payoffs.

Game 6: Prisoner's Dilemma

	t_1	t_2
s_1	8, 8	0, <u>10</u>
s_2	<u>10</u> , 0	<u>2</u> , <u>2</u>

- The DS solution is: (s_2, t_2) .
- How does this solution fare according to our desiderata?
 - It exists for this game, and so satisfies desideratum 2) as well as 1) and 3).
 - Does it seem reasonable to you? (desideratum 5).
 - The solution's payoff sum is $2+2 = 4$, only 25% of the maximum sum $8+8 = 16$, so the solution is inefficient in the transferrable utility sense, and it is also Pareto dominated, so either way, desideratum 4) fails.
 - It predicts well in some settings, as we will see later in the course. So the last desideratum is partially satisfied.

2.3 IDDS

Recall the process illustrated earlier of eliminating strictly dominated strategies, then iterating in the reduced game. If, at the end of this process, only a single profile remains, then it is called the **iterated deletion of dominated strategies IDDS** solution of that game, and the game is said to be dominance solvable.

- IDDS satisfies desideratum 1) — it is always consistent with rationality, and the belief that others are rational. More on this point later.
- we already know from earlier examples that desideratum 2) fails – the IDDS procedure does not always produce a single profile that we can deem to be the solution of the game.

- Apropos 3) [uniqueness], there is a proposition due to Zermelo \sim 1890. Almost every extensive form game (EFG) of perfect information has a unique IDDS solution.
 - The proof idea: Use BI to find the optimal action for each player at each decision node and take expectations across nature moves.
 - Each decision step of BI in the EFG involves eliminating a subset of dominated strategies in the NFG, namely, those strategies involving inferior actions at the contingency (player node or info set) considered at that step.
 - The “almost” in the proposition comes from the fact that there may exist “ties” in the relevant player’s payoffs, in which case the BI solution may not be unique.
- Technical points
 - What if there are several different strategies that are dominated at some iteration — does it matter which you eliminate first? It turns out that the answer is no, it doesn’t matter. See [cite].
 - What if the IDDS procedure halts with some reduced game with at least two players having at least two remaining pure strategies — did we waste our time? The answer is again no; the real action is in the reduced game.
 - * When multiple such profiles remain, they can be refined slightly and considered to be a set of equilibria.
 - * In the reduced game, are any remaining strategies never-best-responses? If so, throw them out to reduce the profile set. Iterate that procedure until no more never-best-replies can be found.
 - * The remaining profiles are called rationalizable equilibria.
 - * They satisfy Des 1) in the strong sense that rationality is “common knowledge.”

- * That last piece of jargon means that each player is rational, and believes that all other players are rational, and believes that all players believe that all players are rational, etc, etc.

See [cite] for a fuller explanation.

- Caveats:

- IDDS may not be feasible. For example, chess is a finite game of perfect information, complete information (win>draw>loss), and perfect recall. Try BI? Chess has an astronomical number of branches. Deep Blue: compute x moves ahead and estimate expected payoffs. This is a “brute force” approach that remains far from IDDS.
- IDDS may not reasonable or empirically valid in certain cases. For example, consider the Centipede game, invented by Robert Rosenthal. It is a finite EFG of complete and perfect information, as in Figure 2.
- In this game there are 2 piles. One pile has two coins, the other zero, in period 1. Two actions: take a pile, (leaving the other pile for the other player and ending the game), or push the two piles across the table to the other player in which case one coin is added to each pile. The game continues for 100 periods or until one player takes a pile.




Figure 2: caption goes here.

- Empirical results: Starting with McKelvey and Palfrey (1992, *Econometrica*), NE play is rarely observed. Explanation is still somewhat controversial, as discussed in the last section of this chapter.

3 Nash equilibrium: Pure Strategies

Nash equilibrium (NE) is the leading solution concept for NFGs. In brief, it is a profile where every player is making a best response to the other players.

3.1 Equal payoff property

- Recall that pure strategy $s_i \in B_i(s_{-i})$, i.e., s_i is a best-response to opponents' profile s_{-i} , if:

$$f_i(s_i, s_{-i}) \geq f_i(s'_i, s_{-i}) \quad \forall s'_i \in S_i. \quad (5)$$

- If this holds for several different pure strategies $s_i \in B_i(s_{-i})$, then they have the same (maximal) payoff, call it m , against s_{-i} . Any weighted average of numbers equal to m is itself equal to m . Therefore, if x_i is a mixture involving only pure strategies in $B_i(s_{-i})$, then $x_i \in B_i(s_{-i})$.

3.2 Definition and existence

A profile $s^* = (s_1^*, \dots, s_n^*)$ of pure strategies in an n -player game is a NE if

$$s_i^* \in B(s_{-i}^*), i = 1, \dots, n. \quad (6)$$

Likewise, a profile $x^* = (x_1^*, \dots, x_n^*)$ of (possibly) mixed strategies is a NE if

$$x_i^* \in B(x_{-i}^*), i = 1, \dots, n. \quad (7)$$

Theorem. John Nash (1951). Let the strategy set for player i , $S_i \subset \Re^m$ be convex and compact (which implies the set is closed and bounded), the payoff function f_i be

continuous in $S = S_1 \times \dots \times S_n$ and, for the restricted payoff function $f_i(\cdot, s_{-i})$ be quasi-concave $i = 1, \dots, n$. Then the game has a Nash equilibrium (not necessarily unique).

Proof sketch:

$$B(s) = (B_1(s_{-1}), \dots, B_n(s_{-n})) \quad (8)$$

defines the best-response correspondence. Can verify that B satisfies the assumptions of Kakutani's fixed point theorem, and therefore has a fixed point. But a fixed point $s^* \in B(s^*)$ is, by definition, a NE.

- Even if $f_i(\cdot, s_{-i})$ is not continuous in $S = S_1 \times \dots \times S_I$ or quasi-concave in $i = 1, \dots, n$, there still may be a NE; see for example pg. 253 MCWG.
- But are the strategy sets convex and compact? The sets of pure strategies are not. In the next section, we'll see that the sets of mixed strategies are, and so we have a
- *Corollary* to this theorem: Every finite normal form game (NFG) has a NE, possibly in mixed strategies.
- Unlike the previous solution concepts, NE does very well on desideratum 2) [existence], in light of Nash's theorem and its corollary.
- It also does well in terms of 5) [reasonable] and 1) [rational] because if a profile is not a NE, then at least one player is not choosing a best-response (not acting rationally).

Here is the first part of a recipe for finding DS, IDDS and pure NE solutions of NFGs. (The second part is for mixed NE.)

0. If given an EFG, write out all players' pure strategy sets (remember: complete contingency plans!) and payoff functions to get the NFG. Collapse identical rows (or columns) to get the R[educed]NFG.

1. Eliminate strictly dominated strategies, including those dominated by mixtures, and write down the reduced NFG. If only one profile remains, it is the DS solution.
2. Repeat step 3) as many times as possible. If only one profile remains, it is IDDS solution.
3. Find each player's pure strategy best responses to each profile of other players' pure strategies. [In bimatrix games there is a handy way to do this by underlining the corresponding payoffs]. Inspect for mutual best responses, i.e., the pure strategy NE.

Let's try out the recipe on

Game 7: CoordinationB

	<i>L</i>	<i>R</i>
<i>U</i>	2, <u>2</u>	<u>2</u> ,2
<i>D</i>	<u>3</u> , <u>1</u>	0,0

1. There are no strictly dominated strategies, although *R* is weakly dominated by *L* for Column player.
2. Moot.
3. Underlining the BR's we find two pure strategy NE: $\{(U, R), (D, L)\}$.

It is hard to argue that Column would play the weakly dominated strategy *R* so it is tempting to eliminate (U, R) . However, Column earns a payoff of 2 at this NE and only 1 at (D, L) . You can't eliminate a weakly dominated strategy without additional criteria or "refinements" (to be discussed later).

loose comment, find a better place for it:

When searching for strictly dominated strategies to eliminate, you should focus on never-best-responses. They may be dominated by pure or mixed strategies

4 Nash Equilibrium: Mixed

The corollary establishing the existence of NE in finite games follows from checking the “mixed extension.” It has strategy sets $\Delta(S_i)$ for player i , and uses the expected payoff as the payoff function.

That is, the strategy sets are simplexes as we have seen, and so convex and compact. We now will see that the payoff functions are linear, hence continuous and quasi-concave.

4.1 Payoffs for mixed strategies.

The payoff function f_i for each player (sometimes denoted u_i or π_i) extends linearly to cover mixed strategies.

- Suppose player i 's pure strategy set is $S_i = \{t_1, \dots, t_n\}$, and she chooses mixture

$$\sigma_i = p_1 t_1 + \dots + p_n t_n.$$

- Player i 's (expected) payoff when the other players use pure strategy profile s_{-i} is

$$f_i(\sigma_i, s_{-i}) = \sum_{k=1}^n p_k f_i(t_k, s_{-i}). \quad (9)$$

- What if some or all of the other players use mixed strategies? Say the profile is

$$\sigma_{-i} = q_1 s_{-i}^1 + \dots + q_m s_{-i}^m.$$

- Then i 's (expected) payoff when playing own pure strategy t_k is

$$f_i(t_k, \sigma_{-i}) = \sum_{j=1}^m q_j f_i(t_k, s_{-i}^j). \quad (10)$$

- If all players are using mixed strategies, i 's (expected) payoff is obtained by substituting equation (10) into (9).

- Later we will see [in evgame chapter] that there are nice compact ways to express such payoffs using matrix algebra.

The last step in recipe, once you have finished the IDDS process and checked for pure NE, is

4. In the remaining strategy profiles, consider each possible combination of subsets of two or more strategies for each player. Set up the system of equations that, for each player, equates her payoffs with the subset of her strategies. If there is solution for which the probabilities are non-negative and sum to 1.0 for each player, then it is a mixed NE.

For example, in Game 7, both pure strategies remain for both players. Thus there is only one relevant subset combination: both strategies for both players.

- Call Row's mixed strategy x where U is chosen with probability p . Thus, $x = pU + (1 - p)D$.
- Call Column's mixed strategy y where L is chosen with probability q . Thus, $y = qL + (1 - q)R$.
- The mixed NE is the simultaneous solution to the set of equations:

$$f_1(U, y) = f_1(D, y) \tag{11}$$

$$f_2(x, L) = f_2(x, R)$$

if the solution satisfies $0 \leq x, y \leq 1$, then it is a mixed NE (and completely mixed if $0 < x, y < 1$).

- If you think about it, there is something peculiar about these two equations.
 - Row's mixing probability p is determined by Column's indifference condition, i.e., Column's payoff function not his own payoffs. Likewise, Column's mixing

probability seems independent of her own payoffs, but directly dependent on Row's payoffs.

- This seems counter-intuitive at first. As we will argue later in more detail, this peculiarity comes from the nature of the best-response correspondence. If $q < q^*$ then pure strategy U is the unique best-response. If $q > q^*$ then pure strategy D is the unique best-response. Only if $q = q^*$ row is indifferent between U and D , and therefore willing to mix U and D .
- It may be worth mentioning that we need cardinal utility when analyzing mixed strategies. Ordinal utility is enough when looking for best responses among pure strategies, and finding pure NE. [[Perhaps move this remark to the previous chapter]]
- Returning to the example, Row player solves

$$\begin{aligned}
 f_1(U, y) &= f_1(D, y). \text{ Since} \\
 f_1(U, y) &= 2q + 2(1 - q) = 2, \text{ and} \\
 f_1(D, y) &= 3q + 0(1 - q) = 3q, \text{ we see that} \\
 f_1(U, y) &= f_1(D, y) \Rightarrow q^* = \frac{2}{3}.
 \end{aligned}$$

- Thus, if Column chooses L with probability $\frac{2}{3}$ and R with probability $\frac{1}{3}$ then Row is indifferent between U and D , and therefore willing to mix U and D . Note the expected payoff of D if $q = q^* = \frac{2}{3}$ is: $\frac{2}{3}(3) + \frac{1}{3}(0) = 2$, which is the payoff of U regardless of q . If $q < q^*$ then pure strategy U is a best-response.
- Column player then solves

$$\begin{aligned}
 f_2(x, L) &= f_2(x, R) \\
 f_2(x, L) &= 2p + 1(1 - p) = 1 + p \\
 f_2(x, R) &= 2p + 0(1 - p) = 2p \\
 f_2(x, L) &= f_2(x, R) \Rightarrow p^* = 1
 \end{aligned}$$

- Thus, Column is indifferent between pure strategies only if $p^* = 1$ since R is weakly dominated.

- Therefore we have a mixed NE: (x, y) , where $x = U$ and $y = \frac{2}{3}L + \frac{1}{3}R$.

Let's apply the recipe to another coordination game

Game 8: CoordinationC

		$(q) \quad (1 - q)$	
		L	R
(p)	U	$\underline{2}, \underline{5}$	$1, 0$
$(1 - p)$	D	$0, 1$	$\underline{5}, \underline{2}$

Again, we can go right to step 3, and find the best responses.

- If $q = 1$ then U (i.e., $p = 1$) is BR
- If $q = 0$ then D (i.e., $p = 0$) is BR
- If $p = 1$ then L (i.e., $q = 1$) is BR
- If $p = 0$ then R (i.e., $q = 0$) is BR

There are two pure NE = $\{(U, L), (D, R)\}$. Clearly Row prefers (D, R) while Column prefers (U, L) .

- Mixed NE. Player 1 solves:

$$\begin{aligned}
 f_1(U, y) &= f_1(D, y) \\
 2q + 1(1 - q) &= 0q + 5(1 - q) \\
 2q &= 4(1 - q) \\
 6q &= 4 \\
 q^* &= \frac{2}{3}
 \end{aligned} \tag{12}$$

– Note: if $q < \frac{2}{3}$, D (i.e., $p = 0$) is BR; if $q > \frac{2}{3}$, U (i.e., $p = 1$) is BR.

- Player 2 solves:

$$\begin{aligned}
 f_2(x, L) &= f_2(x, R) \\
 5p + 1(1 - p) &= 0p + 2(1 - p) \\
 5p &= 1 - p \\
 6p &= 1 \\
 p^* &= \frac{1}{6}
 \end{aligned} \tag{13}$$

– Note: if $p < \frac{1}{6}$, R (i.e., $q = 0$) is BR; if $p > \frac{1}{6}$, L (i.e., $q = 1$) is BR.

- The mixed NE is where player 1 chooses U with probability $p = \frac{1}{6}$ and D with probability $\frac{5}{6}$ while player 2 chooses L with probability $\frac{2}{3}$ and R with probability $\frac{1}{3}$.

4.2 Best-response diagram

- We can draw the best-response correspondences and find the mixed NE as their intersection. (Not best-response *functions*, since the graphs include vertical lines.)
- We have, $f_2(x, L) - f_2(x, R) = 1 + 4p - (2 - 2p) > 0$ if $6p > 1$ or $p > p^* = \frac{1}{6}$.
 - Thus, for any $p > p^* = \frac{1}{6}$ pure strategy L ($q = 1$) is a best-response and for any $p < p^* = \frac{1}{6}$ pure strategy R ($q = 0$) is a best-response.
 - Any $q \in [0, 1]$ is a best-response if $p = p^* = \frac{1}{6}$.

$$p^*(q) = BR_1 = \left\{ \begin{array}{ll} 0 & \text{if } q < \frac{2}{3} \\ 1 & \text{if } q > \frac{2}{3} \\ [0, 1] & \text{if } q = \frac{2}{3} \end{array} \right\}$$

- Similarly,

$$q^*(p) = BR_2 = \begin{cases} 0 & \text{if } p < \frac{1}{6} \\ 1 & \text{if } p > \frac{1}{6} \\ [0, 1] & \text{if } p = \frac{1}{6} \end{cases}$$




Figure 3: Best response correspondences for Game 8. The pure NE are the corner points where the two correspondences intersect, and the mixed NE is the interior intersection.

- The NFG looks like it could be symmetric, but the mixing probabilities are different.
 - Players 1 and 2 have similar payoffs at the two pure NE (5 or 2 for a total of 7) and at the off diagonal elements (0 and 1).
 - Later we will define a symmetric game and show that this game is not symmetric, thus the mixing probabilities are different.
- Each player's payoffs at the mixed NE are equal across the pure strategies that she mixes: $(x = \frac{1}{6}U + \frac{5}{6}D, y = \frac{2}{3}L + \frac{1}{3}R)$.

– Player 1:

$$\begin{aligned} E[\pi_1] &= \frac{1}{6} \left[\frac{2}{3}(2) + \frac{1}{3}(1) \right] + \frac{5}{6} \left[\frac{2}{3}(0) + \frac{1}{3}(5) \right] \\ E[\pi_1] &= \frac{1}{6} \left[\frac{5}{3} \right] + \frac{5}{6} \left[\frac{5}{3} \right] = \frac{5}{3} \end{aligned}$$

– Player 2:

$$\begin{aligned} E[\pi_2] &= \frac{2}{3} \left[\frac{1}{6}(5) + \frac{5}{6}(1) \right] + \frac{1}{3} \left[\frac{1}{6}(0) + \frac{5}{6}(2) \right] \\ E[\pi_2] &= \frac{2}{3} \left[\frac{10}{6} \right] + \frac{1}{3} \left[\frac{10}{6} \right] = \frac{5}{3} \end{aligned}$$

* Note: $\frac{5}{3} < 2 < 5$, so mixed NE payoff is lower than either pure NE payoff.

* Note: $\frac{5}{3}$ is an *expectation* and not a *realization* outcome.

· If the NFG came from an EFG with Nature moves, then the NFG payoffs are already expectations, and the mixed NE payoffs are expected values of expected values.

4.3 Mixed NE as maximum expected payoff

Continuing with Game 8,

- Maximizing expected payoff under uncertainty about the other players mixed strategy offers an alternative method of achieving the same mixed NE.

- Again, let player 1 (row player) choose U with probability p and player 2 (column player) choose L with probability q .
- Player 1 chooses p to maximize expected payoff as a function of q .

$$\begin{aligned} \max_p E(\pi_1) &= q[2p + 0(1 - p)] + (1 - q)[1p + 5(1 - p)] \\ &= 2pq + p - pq + 5(1 - q)(1 - p) \\ &= pq + p + 5 - 5q - 5p + 5pq \\ &= 6pq - 4p - 5q + 5 \end{aligned} \tag{14}$$

* The first order condition is

$$\begin{aligned}\frac{\partial E(\pi_1)}{\partial p} &= 6q - 4 = 0 \\ q^* &= \frac{2}{3}\end{aligned}\tag{15}$$

* Note:

$$\begin{aligned}\frac{\partial E(\pi_1)}{\partial p} &< 0 \text{ if } q < \frac{2}{3} \implies p^* = 0 \\ \frac{\partial E(\pi_1)}{\partial p} &> 0 \text{ if } q > \frac{2}{3} \implies p^* = 1.\end{aligned}$$

* This is the exact same result we obtained in equations (13) and (12). Note the SOC are zero and don't confirm (or contradict) a maximum.

– Player 2 maximizes expected payoff by choosing q .

$$\begin{aligned}\max_q E(\pi_2) &= p[5q + 0(1 - q)] + (1 - p)[1q + 2(1 - q)] \\ &= 5pq + q - pq + 2(1 - p)(1 - q) \\ &= 4pq + q + 2 - 2p - 2q + 2pq \\ &= 6pq - q - 2p + 2\end{aligned}\tag{16}$$

* The first order condition is

$$\begin{aligned}\frac{\partial E(\pi_2)}{\partial q} &= 6p - 1 = 0 \\ p^* &= \frac{1}{6}\end{aligned}\tag{17}$$

* Note:

$$\begin{aligned}\frac{\partial E(\pi_2)}{\partial q} &< 0 \text{ if } p < \frac{1}{6} \implies q^* = 0 \\ \frac{\partial E(\pi_2)}{\partial q} &> 0 \text{ if } p > \frac{1}{6} \implies q^* = 1.\end{aligned}$$

5 Refinements, QRE and CE

Many game theorists feel that the most important weakness of NE as a solution concept is the lack of uniqueness. Many games have more than one NE. In that case, which of them should be taken seriously as predictions, and which should not? How should we refine the set of NE so that we get closer to a unique prediction?

John Harsanyi pursued this question relentlessly, and he and Selten wrote a book on the subject. Many top game theorists joined in the quest for NE refinements, especially in the 1980s and 1990s, and came up with dozens of ingenious ideas. In the end, however, the quest failed, and there is still no generally accepted way to select a single NE as “the” appropriate prediction. Yet the quest generated many helpful ideas, and we will survey a few that seem to withstand the test of time.

5.1 Payoff dominance vs risk dominance.

1. Payoff dominance and risk-dominance are Harsanyi and Selten’s top two refinement ideas.
2. If a NE delivers lower payoffs to all players than another NE, then it seems less attractive and perhaps should be taken less seriously. So one refinement is *Pareto dominance* – NE that are not Pareto dominated by other NE.
3. Risk dominance means robust to strategic uncertainty.
 - A NE is risk dominant if, in some appropriate sense, players are not very vulnerable to reduced payoffs when other players occasionally deviate from their NE strategies.
 - There are several ways to formalize the “appropriate sense.”
 - The best way (to our knowledge, in terms of explaining regularities observed

in the lab) is in terms of a wide (or deep) basin of attraction. This will be defined in the chapter on evolutionary games.

- Later in this section we will see that Trembling-hand perfection (THPNE) is a refinement closely related to risk dominance.

4. The next chapter will explain Subgame perfection (SPNE) and its connection to THPNE.

Consider the following symmetric game:

Game 9: Coordination

	L	R
T	5, 5	-1, 4
B	4, -1	3, 3

- There are two pure NE: (T, L) and (B, R)
- (T, L) is payoff dominant – it has a higher payoff than the NE (B, R) for both players .
- On the other hand, the NE (B, R) is risk dominant – players are not hurt as much (or at all!) by deviations from this equilibrium.
- We can ask how robust each of the pure NE are to deviations.

– First consider (B, R) .

- * Suppose that player 2 is intending to choose R , but deviates with probability ϵ and chooses L .
- * The payoff advantage for player 1 from choosing B is now:

$$\begin{aligned}
 E\pi_1(B) - E\pi_1(T) &= [4\epsilon + 3(1 - \epsilon)] - [5\epsilon - 1(1 - \epsilon)] \\
 &= [3\epsilon + 3] - [6\epsilon - 1] \\
 &= 4 - 3\epsilon > 0 \quad \forall \epsilon \in [0, 1]
 \end{aligned} \tag{18}$$

- * So, for any $\epsilon > 0$ the expected payoff from playing B is greater than T , thus (B, R) is robust any deviation by player 2.
- * The same math shows that (B, R) is robust any deviation by player 1.
- Compare this with the other pure NE: Player 2 deviates from the NE (T, L) and chooses R with probability $\epsilon \in [0, 1]$.
- * The payoff advantage for player 1 from choosing T is now:

$$\begin{aligned}
E\pi_1(T) - E\pi_1(B) &= [5(1 - \epsilon) - 1\epsilon] - [4(1 - \epsilon) + 3\epsilon] \\
&= [5 - 6\epsilon] - [4 - \epsilon] \\
&= 1 - 5\epsilon
\end{aligned} \tag{19}$$

- * So, for any $\epsilon < .2$ the expected payoff from playing T is greater than B , thus (T, L) is robust of to an 20% tremble rate by player 2.
 - Thus, (B, R) is risk-dominant since it remains a best-response for a larger degree of strategic uncertainty.
- One definition of risk dominance states that NE players are still making their best responses as long as opponents probability of a tremble is less than 0.5 (for 2 pure strategies).
 - Before moving on, we note that the number of NE is usually odd. Since we have two pure NE, it is likely that there is also a mixed NE for the game we are looking at.

- Using the expected payoff method in equations (16) through (15) we get:

$$\begin{aligned}
\max_p E(\pi_1) &= p[5q + -1(1 - q)] + (1 - p)[4q + 3(1 - q)] \\
&= p[6q - 1] + (1 - p)[q + 3] \\
&= 5pq - p + q - 3p + 3 \\
\frac{\partial E(\pi_1)}{\partial p} &= 5q - 4 = 0 \\
q^* &= \frac{4}{5}
\end{aligned} \tag{20}$$

- Since this game is symmetric $p^* = \frac{4}{5}$ as well.

- The mixed NE $(x, y)(x = \frac{4}{5}T + \frac{1}{5}B, y = \frac{4}{5}L + \frac{1}{5}R)$.

- The chapter on evolutionary games will show that it is no coincidence that the mixture probability $\frac{1}{5} = 0.2$, the maximum deviation rate.

5.2 Correlated equilibrium

In some cases, instead of narrowing down the set of NE, we want to expand the solution set beyond NE. Besides rationalizable equilibria (discussed earlier), the leading such solution concept is called correlated equilibrium (CE). It is often relevant in computer science applications of game theory, and occasionally elsewhere.

- A mixed strategy NE assumes that the mixing probabilities p and q are statistically independent.
- For example, in a 2x2 bimatrix game, the “probability matrix” showing the weight of the utility outcomes in the expected payoffs is:

		Column (2)	
Row (1)		$t_1(q)$	$t_2(1 - q)$
p	s_1	$pq(\pi_1, \pi_2)$	$p(1 - q)(\pi_1, \pi_2)$
$(1 - p)$	s_2	$(1 - p)(q)(\pi_1, \pi_2)$	$(1 - p)(1 - q)(\pi_1, \pi_2)$

- In a correlated equilibrium, the cell probabilities could differ from above, while keeping the column and row sum probabilities constant.
- In a correlated equilibrium the probability of s_1 is not necessarily p , but may depend on the realization of q , that is we can have $\text{prob}(p|q) \neq p$ and $\text{prob}(q|p) \neq q$.

Recall, for example, game 4: Battle of the Sexes.

- It has 3 NE: (s_1, t_2) , interpreted as both players going to the opera; (s_2, t_1) interpreted as both going to the boxing match; and $(.5s_1 + .5s_2, .5t_1 + .5t_2)$, interpreted as random choice.
- The (expected) payoff sum is 60 for the two pure NE, but one player is happier than the other.
- The mixed NE has expected payoff sum only 30, and it is Pareto dominated.
- It is a correlated equilibrium if the players flip just one fair coin, and play the first pure NE if Heads and play the second pure NE if Tails.
- Here $p = q = .5$ still, but the realizations are perfectly negatively correlated.
- See (cite) for more on CE.

6 Behavioral Considerations.

Centipede. The Centipede game is famous because it pushes BI (and IDDS) to the breaking point. As noted, lab experiments since Palfrey and xx have found that players don't tend to grab on the first opportunity, as predicted in NE (or IDDS). They let it ride for a while, but not to the very end.

Here is the basic issue as we see it. The logic for grabbing immediately is close to complete common knowledge of rationality – you have to [believe that your opponent

believes]¹⁰⁰ that both are rational. If there is a seed of doubt in your mind, it may be worth letting it ride for a while, since little is lost and there is much to be gained. But as the game goes on, there is more and more to be lost, and less and less to be gained. So the observed behavior actually makes good sense.

We will return to such matters in the last section of Chapter 6, when we discuss the finitely repeated Prisoners dilemma.

Risk vs payoff dominance. Laboratory experiments generally give stronger support to risk dominance than to payoff dominance. For example, in weakest link games...

But payoff dominance also has some support. Friedman (1996) looked at games similar to Game 9 (CoordinationD). Observed play was best explained as if the tremble rate ϵ_{PD} for the Pareto dominant NE was less than the tremble rate ϵ_{RD} for the risk dominant NE. The interpretation is that the RDNE is indeed less dangerous than the PDNE, but not as risky as one might think. Again, this will be discussed further in Chapter 6.

Quantal Response Equilibrium (QRE). Suppose tremble rates are larger when less is at stake. E.g., consider logit or probit distribution for random utility. Then suppose everyone is aware of the common tremble rate. Derive QRE, discovered independently by Palfrey and xx, and by James Friedman and yy. Write down formula, discuss literature, including recent Selten et al article comparing to other bounded rational perturbations of NE.

biblio Harsanyi and Selten (1988). “A General Theory of Equilibrium Selection in Games”

Holt and Palfrey book on QRE.

Friedman (1996) Economic Journal.